

## SELF-ASSESSMENT



# Agent Evaluation Scorecard

A self-assessment for teams building LLM agents. Mark where your team is today — honestly. The point is to find the gaps, not to feel good.

For each dimension below, mark where your team is today. Score each from **0** to **3** using the scale, then add it up out of 24.

**0**

### Absent

We don't do this at all.

**1**

### Ad-hoc

Sometimes, manually, no system.

**2**

### Established

A defined process. It runs.

**3**

### Mature

Automated, continuous, feeds back into development.

## 1 Defined success criteria

Can anyone on your team write down in one paragraph what "good" looks like? Not "it answers correctly" — concrete, measurable, agreed across engineering, product, and the people who actually use it.

☐ 0 ☐ 1 ☐ 2 ☐ 3

## 2 Evaluation dataset

A curated set of inputs covering typical, ambiguous, edge, and known-failure cases. It is version-controlled. It grows.

☐ 0 ☐ 1 ☐ 2 ☐ 3

## 3 Automated evaluation runs

You can run your evaluation suite on demand and on every meaningful change — prompt, model, tool, code. Results are recorded.

☐ 0 ☐ 1 ☐ 2 ☐ 3

## 4 Trace visibility

For any failure or surprising output, you can see the full execution: prompts, tool calls, intermediate steps, model responses, latency, cost. Not just the final answer.

☐ 0 ☐ 1 ☐ 2 ☐ 3

## 5 Reproducible failures

When a user reports a problem or an evaluation fails, you can re-run the exact case and reliably see the same failure. You don't lose bugs to "couldn't reproduce."

☐ 0 ☐ 1 ☐ 2 ☐ 3

## 6 Version comparison

You can compare two versions of your agent — different prompts, models, tool descriptions — side by side on the same dataset. You see which is better, and where.

☐ 0 ☐ 1 ☐ 2 ☐ 3

---

## 7 Production monitoring

You see how the agent behaves in production with real users, not just on your curated dataset. Production failures flow back into the evaluation dataset.

☐ 0 ☐ 1 ☐ 2 ☐ 3

---

## 8 Ownership of quality

There is a named person or team accountable for agent quality. They have the time, tools, and authority to act on what evaluation reveals.

☐ 0 ☐ 1 ☐ 2 ☐ 3

## How to read your result — total / 24

0–8

**The “seems to work” zone.** Most teams are here. The risk is that you only find out what’s broken when a user does. Start with dimensions 1, 2, and 4.

9–16

**Some structure, not yet a system.** Typical gap: you have a dataset and traces, but no automated runs and no production feedback loop. Connect what you already have.

17–24

**A real evaluation discipline.** The remaining work is depth — better metrics, harder edge cases, faster iteration. The structure is in place.

### A NOTE ON DIMENSION 1

This is the one most teams skip and most regret skipping. **“What does good look like?”** is not an engineering question — it requires product, domain experts, and the people who will actually use the output. If you can’t answer it, no amount of tooling setup will help.