

SELF-ASSESSMENT



Agent Evaluation Scorecard

A self-assessment for teams building LLM agents. Mark where your team is today. Use the gaps to decide what to improve next.

For each dimension below, mark where your team is today. Score each from **0** to **3** using the scale, then add it up out of 24.

0

Absent

We don't do this at all.

1

Ad-hoc

Sometimes, manually, no system.

2

Established

A defined process. It runs.

3

Mature

Applied consistently, reviewed, and used to improve development.

1 Defined success criteria

You define successful behavior before building or changing the agent. Concrete examples and measurable criteria reflect the needs of engineering, product, and the people who use it.

☐ 0 ☐ 1 ☐ 2 ☐ 3

2 Evaluation dataset

A versioned set of tasks and success criteria covers typical cases, ambiguity, and known failures. Keep successful cases too. Expand coverage as you learn; a small starter set is not proof of production reliability.

☐ 0 ☐ 1 ☐ 2 ☐ 3

3 Automated evaluation runs

You run evaluations on demand and on meaningful changes to prompts, models, tools, or code. Results are recorded. Grading rules are checked against human judgment where appropriate.

☐ 0 ☐ 1 ☐ 2 ☐ 3

4 Trace visibility

You can inspect recorded prompts, tool calls, results, model responses, latency, and cost to investigate failures. You know what your instrumentation captures and where visibility is incomplete.

☐ 0 ☐ 1 ☐ 2 ☐ 3

5 Repeatable investigation

You capture inputs, versions, and relevant environment state so you can rerun a failing case. For variable behavior, repeat trials to measure how often it fails. One passing rerun does not establish a fix.

☐ 0 ☐ 1 ☐ 2 ☐ 3

6 Version comparison

You compare versions on the same cases and grading criteria, using repeated trials when needed. Inspect regressions as well as averages. Check improvements on additional cases before generalizing.

☐ 0 ☐ 1 ☐ 2 ☐ 3

7 Production monitoring

You see how the agent behaves in production with real users, not just on your curated dataset. Production failures flow back into the evaluation dataset.

☐ 0 ☐ 1 ☐ 2 ☐ 3

8 Ownership of quality

There is a named person or team accountable for agent quality. They have the time, tools, and authority to act on what evaluation reveals.

☐ 0 ☐ 1 ☐ 2 ☐ 3

How to read your result - total / 24

These bands are a discussion guide, not a validated benchmark or release gate. Review each dimension: a high total can hide a critical gap.

0-8

Build the foundations. Start with explicit success criteria, a small set of representative cases, and enough trace visibility to investigate failures.

9-16

Connect the pieces. Review your lowest-scoring dimensions. Link evaluation runs, failure investigation, and production feedback into a regular process.

17-24

Test the depth. Challenge dataset coverage, grader quality, and consistency across runs. A strong self-assessment score does not establish production reliability.

START WITH EXPECTATIONS

Define what success looks like before you build, and make those expectations something you can repeatedly test. Start with a few real tasks and clear criteria. Grow the checks as you learn from users and failures.